Business Models for Agent Reliability Testing

Every realistic way to charge for evaluating and regression-testing AI agents, for discussion with my cofounder

1. The core idea in one paragraph

We find out whether a company's AI agent picks the right tool, avoids wasted calls, calls tools in the right order, and respects business rules. An agent changes every time the prompt, tool list, model or code changes, so a one-time check goes stale quickly. That is why **regression testing** (re-running the same checks after every change) is the natural recurring product: it turns a one-off audit into repeat revenue.

Every model below is a different answer to two questions: **when** do we test (the trigger), and **how** do we get paid (the pricing unit).

2. The regression-testing triggers (your idea, fully listed)

Each row is a moment when an agent's behavior can silently change. Each can be a product on its own.

Trigger	When it runs	What it catches	Who feels the pain
Every production deploy	Automatically on each release to production (CI/CD step)	New code broke tool choice, arguments or rules	Engineering teams that ship often
Every version / release	Before each tagged version goes out (release gate)	Behavior differences between version N and N-1	Teams with a release process or QA sign-off
Every pull request	On each code change, before merge	Problems caught earliest, cheapest to fix	Larger teams with CI already in place
Prompt change	Whenever the system prompt is edited	A wording tweak that fixes one case and breaks three	Everyone; prompts change most often
Tool change	Tool added, removed, renamed or its description edited	Agent now confuses two similar tools	Teams growing their tool list
Model change or upgrade	Switching vendor or version (e.g. a new model release)	Same prompt behaves differently on a new model	Any team on a hosted model, which updates under them
Scheduled (nightly / weekly)	On a timer, even with no code change	Silent model drift, data changes, API changes	Teams with no formal release process
Live production sampling	Continuous scoring of a sample of real traffic	Failures real users hit that tests never covered	Teams with real traffic and cost or safety exposure
Incident-triggered	After a bad outcome or customer complaint	Root cause, plus a new test so it never repeats	Teams after a visible failure
Business-rule change	Policy changes (e.g. new refund limit)	Agent still follows the old rule	Regulated or policy-heavy businesses

The flywheel: every failure found in production becomes a new test case in that client's golden set. The longer a client stays, the better their test suite and the harder we are to replace. That switching cost is the real moat for the recurring models.

3. Service business models

Humans plus our engine. Fast to start, no product to finish first, and they teach us what to build.

S1. One-time audit

What it is	Fixed-scope pilot: collect tools and rules, build the test set, run the checks, deliver a report with fixes.
Who buys	Teams with an agent in production or close to launch.
How it's priced	Fixed fee per audit.
Upside	Easiest first sale. Produces case studies and real failure patterns.
Watch out for	One-off revenue. Does not repeat unless paired with S2 to S5.

S2. Pre-launch readiness review (launch gate)

What it is	A go / no-go assessment before an agent launches or expands to a new use case.
Who buys	Teams with a launch date, or regulated teams needing sign-off.
How it's priced	Fixed fee per launch.
Upside	Urgency is built in. Buyers have a deadline and budget.
Watch out for	Tied to launch calendars. Lumpy demand.

S3. Per-release regression testing (every version)

What it is	We re-run the client's golden set against each new version and deliver a pass / fail diff versus the previous version.
Who buys	Teams that release on a cadence (monthly, biweekly).
How it's priced	Fee per release, or a bundle of N releases per quarter.
Upside	Repeat revenue with low delivery effort once the test set exists. Directly addresses your 'every version' idea.
Watch out for	Revenue depends on how often the client ships. Clients may push to automate it, which leads to P3.

S4. Model-migration testing

What it is	Before a client switches model vendor or version, we compare old versus new on their tests and report regressions and cost differences.
Who buys	Teams moving to a cheaper or newer model.
How it's priced	Fixed fee per migration.
Upside	Clear ROI story: switching models can save cost, but only if quality holds. Happens often as models update.
Watch out for	Episodic. Needs a mature test set to be credible.

S5. Retainer: ongoing monitoring and re-runs

What it is	Monthly service: scheduled re-runs, a trend report, new test cases from fresh traces, and a standing review call.
Who buys	Teams without a dedicated evaluation person.
How it's priced	Monthly fee, tiered by number of agents or review depth.
Upside	Most predictable revenue for a service. Builds the flywheel.
Watch out for	Does not scale without automation. Your time is the product.

S6. Incident investigation and post-mortem

What it is	After a bad agent outcome, we trace the root cause, reproduce it, and add a regression test.
Who buys	Teams after a customer-facing failure.
How it's priced	Fixed fee per incident, or hourly.
Upside	High urgency and willingness to pay. Strong trust-builder.
Watch out for	Reactive and unpredictable. Needs fast turnaround.

S7. Fix implementation consulting

What it is	We do not just report. We rewrite tool descriptions, prompts and guardrails, then re-test to prove improvement.
Who buys	Teams without time or expertise to act on a report.
How it's priced	Project fee or hourly.
Upside	Higher revenue per client. Closes the loop from finding to fixing.
Watch out for	Most labor-intensive. Risk of becoming a general AI consultancy.

S8. Independent assessment report

What it is	A third-party report the client can show to their own customers, auditors or investors.
Who buys	Regulated areas (health, finance) and companies selling agents to enterprises.
How it's priced	Fixed fee per report, optionally renewed yearly.
Upside	Credibility is the product, which a self-built dashboard cannot provide.
Watch out for	Needs a track record and a clear, defensible method. Avoid claiming formal certification we do not hold.

S9. Training and workshops

What it is	Teach a client's engineers how to build golden sets and evaluate their own agents.
Who buys	Teams that want to build the capability in-house.
How it's priced	Fixed fee per workshop.
Upside	Low delivery cost once prepared. Good lead generation.
Watch out for	Teaches clients to not need us. Best as a bundle with other services.

4. Product (software) business models

The same engine, packaged so clients can use it without us in the loop. Slower to build, scales much better.

P1. Subscription per agent

What it is	Flat monthly price per agent under test, in tiers (more tests, more runs, more history).
Who buys	Teams with one to a few agents.
How it's priced	Monthly or annual, per agent.
Upside	Simple to understand and budget. Predictable revenue.
Watch out for	Heavy users cost more to serve than light users pay.

P2. Usage-based pricing

What it is	Pay per evaluation run, per trace scored, or per test case executed.
Who buys	Teams with spiky or fast-growing volume.
How it's priced	Per run, per trace or per test, often with a free allowance.
Upside	Revenue grows with the customer. Low barrier to start.
Watch out for	Unpredictable bills worry buyers. Compare cheap observability tools.

P3. Automated regression suite in CI/CD (every deploy)

What it is	A pipeline step (such as a GitHub Action) that runs the client's golden set on every deploy and blocks a release if scores drop.
Who buys	Engineering teams with CI/CD.
How it's priced	Subscription plus a run allowance; extra runs charged.
Upside	The strongest recurring product: embedded in their workflow and painful to remove. Automates S3.
Watch out for	Needs the client to maintain a good golden set. Needs reliable, repeatable scoring to avoid false alarms that teams ignore.

P4. Continuous production monitoring with alerts

What it is	Samples live traffic, scores it continuously, and alerts when tool accuracy, wasted calls, loops or rule breaks spike.
Who buys	Teams with real users and cost or safety exposure.
How it's priced	Subscription tiered by traffic volume.
Upside	Catches failures that tests never imagined. Sticky because monitoring runs daily.
Watch out for	Crowded market (observability vendors). Needs a differentiator beyond dashboards.

P5. Self-hosted or on-premise license

What it is	Clients run the engine inside their own environment, so traces never leave their network.
Who buys	Health, finance, government and other privacy-sensitive buyers.
How it's priced	Annual license plus support.
Upside	Removes the biggest sales objection (data leaving the building). Higher contract values.
Watch out for	Support burden and slower sales cycles. Check licensing of any open-source components we build on.

P6. Domain packs (add-on)

What it is	Prebuilt tests and rule templates per industry: healthcare booking, customer support, RAG, finance.
Who buys	Clients in the matching industry.
How it's priced	Add-on subscription or one-time purchase.
Upside	Faster onboarding and a defensible library that grows with every audit.
Watch out for	Build only after seeing the same patterns across several clients. Building too early means guessing.

P7. Dataset and benchmark products

What it is	Curated golden datasets, adversarial test sets or benchmarks that teams license to test their own agents.
Who buys	Teams that want ready-made tests, and agent vendors that want credibility.
How it's priced	License or subscription per dataset.
Upside	Reusable asset with near-zero marginal cost.
Watch out for	Value depends on quality and freshness. Public benchmarks exist, so ours must be more specific or more current.

P8. Outcome-based pricing

What it is	A share of measured savings, such as reduced wasted tool calls and lower model spend, or fewer failed transactions.
Who buys	Teams with significant model or tool-call cost.
How it's priced	Percentage of verified savings, usually with a minimum fee.
Upside	Easy to justify. Aligns our incentive with the client's.
Watch out for	Needs trusted measurement and clean baselines. Slow to prove and negotiate.

5. Hybrid, channel and platform models

H1. Audit-to-SaaS (service first, product later)

What it is	Start with paid audits (S1). Automate whatever we repeat every time into software (P1 to P4). Convert audit clients into the first customers.
Who buys	Early clients who liked the audit.
How it's priced	Audit fee, then subscription.
Upside	Lowest risk path. We build only what clients prove they will pay for. This is the recommended start.
Watch out for	Takes discipline to stop doing manual work and productize it.

H2. Open-core

What it is	Free open-source evaluation engine; paid hosted version, enterprise features, support and domain packs.
Who buys	Developers who adopt for free; companies that pay to scale.
How it's priced	Free core; paid hosted and enterprise tiers.
Upside	Distribution and trust. Developers adopt without sales calls.
Watch out for	Most users never pay. Needs community effort. Hard to switch strategy later.

H3. White-label and partner channel

What it is	Sell the engine or reports to AI agencies and consultancies who resell it under their own brand to their clients.
Who buys	Agencies and system integrators building agents for others.
How it's priced	Wholesale subscription or revenue share.
Upside	One partner brings many clients. Agencies have the trust we lack.
Watch out for	Dependence on partners. Lower margin per client.

H4. API and embedded evaluation

What it is	Agent-building platforms embed our evaluation as a feature for their users.
Who buys	Agent framework and platform vendors.
How it's priced	Usage-based API fees or revenue share.
Upside	Large distribution from a single deal.
Watch out for	Platforms may build it themselves. Long enterprise sales cycle.

H5. Content, benchmarks and lead generation

What it is	Publish public failure-pattern reports and benchmarks on common agent types (for example, tool-choice errors) to attract clients.
Who buys	Not a revenue stream by itself. It feeds every other model.
How it's priced	Free; paid by the clients it brings in.
Upside	Builds credibility with little cost. Our audits become the raw material.
Watch out for	Never publish client data. Aggregate and anonymize, with permission.

6. Side-by-side comparison

Model	Build effort	Revenue type	Scales?	Data privacy friction
S1 One-time audit	Low	One-off	No	Medium
S2 Launch gate	Low	One-off	No	Medium
S3 Per-release testing	Low to medium	Repeat	Somewhat	Medium
S4 Model migration	Low to medium	Episodic	Somewhat	Medium
S5 Monthly retainer	Low	Recurring	No	Medium
S6 Incident work	Low	Episodic	No	High
S7 Fix consulting	Medium	Project	No	High
S8 Assessment report	Medium	Yearly	Somewhat	High
S9 Workshops	Low	One-off	Somewhat	Low
P1 Subscription	High	Recurring	Yes	Medium
P2 Usage-based	High	Variable	Yes	Medium
P3 CI/CD regression	High	Recurring	Yes	Medium
P4 Production monitoring	High	Recurring	Yes	High
P5 Self-hosted license	High	Annual	Yes	Low
P6 Domain packs	High	Add-on	Yes	Low
P7 Datasets	Medium to high	License	Yes	Low
P8 Outcome-based	Medium	Variable	Somewhat	High
H1 Audit-to-SaaS	Staged	Mixed	Yes, over time	Medium
H2 Open-core	High	Mixed	Yes	Low
H3 White-label	Medium	Wholesale	Yes	Medium
H4 Embedded API	High	Usage	Yes	Medium

Ratings are my rough judgment for discussion, not measured data.

7. Suggested path

The point is to sequence these, not pick one forever.

Phase 1: Prove the engine (now). Use VetMedd and a toy agent as test beds. Produce real audit numbers.

Phase 2: Paid pilots. Sell S1 and S2 (audit, launch gate) to two or three teams. Offer S3 (per-release regression) as the natural follow-up. Learn what they repeat.

Phase 3: Recurring. Move the best pilots to S5 and S3 (retainer and per-release). Write down every check we run by hand each time.

Phase 4: Automate. Turn those repeated checks into P3 (CI/CD regression suite) and P4 (monitoring). This is the first real software product.

Phase 5: Expand. Add P5 (self-hosted) for privacy-sensitive buyers, P6 (domain packs) once patterns repeat, and H3 (partner channel) for reach.

8. Open questions for us to decide

- Which first customer type: small teams without evaluation staff, or regulated companies (health, finance)?
- Do we price per release, per month, or per run? Per-release fits our idea, but clients with frequent deploys will want a flat fee.
- How much do we promise to fix versus only report? Fix work pays more but pulls us toward general consulting.
- What is our stance on client data: redacted only, local only, or on-premise from day one?
- What is our real differentiator against observability tools that already ship evaluators: our method, our rule library, or our cross-client pattern knowledge?
- Which part of the engine do we open-source, if any?
- Who builds and who sells? A founder-led service phase needs one of us on client work.

9. Risks to keep in mind

- **Commodity risk:** tracing and generic checks are being bundled into existing tools. Our value must be the golden sets, rules, validated judges and fixes.
- **False alarms:** a regression suite that cries wolf gets switched off. Scoring must be deterministic wherever possible, and LLM judges must be validated against hand-labeled cases.
- **Golden set quality:** every score is only as good as the expected answers. Poor labels mean poor numbers.
- **Client dependence:** a client who ships rarely pays rarely under per-release pricing. Mix in a monthly minimum.
- **Licensing:** check the licenses of any open-source tools we build on before offering them as a hosted product. Get legal advice before launch.
- **Overclaiming:** do not call a report a formal certification unless we hold the standing to do so.