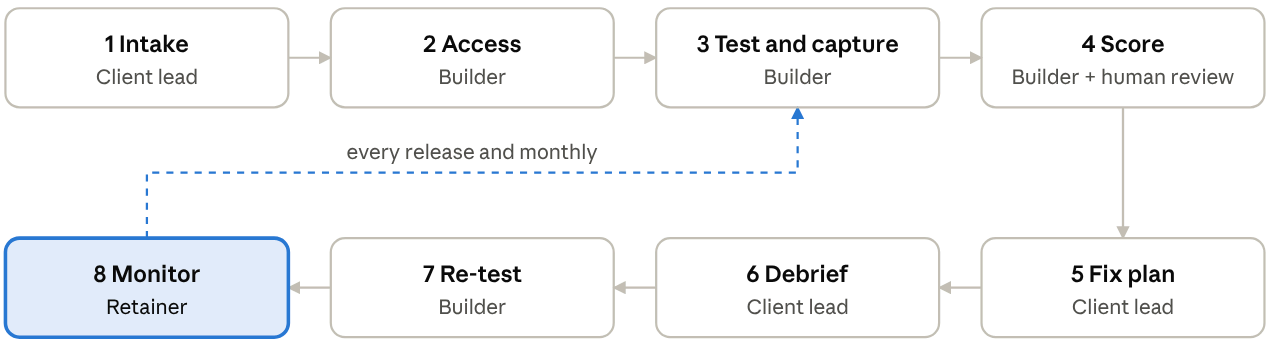
Engagement Loop

Oct 17, 2026 · @Sharkk

Every client goes through the same eight steps: intake, access, test and capture, score, fix plan, debrief, re-test, monitor. Positioning: a complete health check for your AI agent — what's wrong, why, and what to fix first — and continuous monitoring so it stays healthy.

At a glance

Every audit runs eight steps; monitoring repeats the test on each release



engagement loop · 8 steps, monitoring repeats

Steps 1–7 are the one-off audit. Step 8 is the retainer: it re-runs the test on every release, which is where recurring revenue comes from.

Step by step

Each step has one owner and a clear finish line. Client-facing work is the client lead's; anything that touches the agent or code is the builder's.

#	Step	What happens	Owner	Output	Done when
1	Intake	Collect the tool list, business rules and targets on a 30-minute call. Draft the test set (about 50 questions with expected behaviour) and get the client to approve it.	Client lead	Intake doc + approved test set	Client signs off on rules, targets and test set
2	Access	Client picks endpoint, logs, tracing or code access. Send the setup guide, connect, and confirm one test chat comes through.	Builder	Working connection to the agent	One test chat recorded end to end
3	Test and capture	Send every test question three times (for consistency). Pull traces or logs of real behaviour where available.	Builder	Raw results: replies, tool calls, timings	Every test case has a recorded run
4	Score	Rule checks for tool, arguments, necessity, sequence and loops. A judge model checks grounding, task success and rules. A human reviews about 15% of cases before any number is final.	Builder runs, client lead reviews	Results CSV + scorecard vs targets	Spot-checked scores match what a person would mark
5	Fix plan	Turn each failure pattern into a finding: what happens, how often, evidence, root cause, fix, effort. Rank into Now, Next, Later.	Client lead writes, builder confirms causes	Finding cards + roadmap	Every finding has frequency, evidence, cause and fix
6	Debrief	Send the branded report and CSV, then walk the client team through it on a 30-minute readout.	Client lead	Report, CSV, readout call	Client agrees on the top fixes and dates

#	Step	What happens	Owner	Output	Done when
7	Re-test	After the client ships fixes, re-run the same test set and compare before and after.	Builder	Before / after scorecard	Score change reported to the client
8	Monitor	Re-run the test set on every release and monthly, flag regressions, send a short monthly report.	Builder automates, client lead reports	Monthly report + release alerts	Retainer signed and runs automated

Access options

The client chooses how much we see. Start with the least access that answers their question, and let the report show what deeper access would add.

Option	Client gives	Change on their side	We can tell them
A. Endpoint	An API or chat URL and a test account	None	What goes wrong: wrong answers, made-up facts, unsafe or slow replies
B. Log or trace export	Redacted logs of past runs	None	What happened with real users, as far as their logs record it
C. Self-install tracing	Their engineer runs our setup file (about 15 minutes)	One new file, two small edits; fail-safe and reversible	Why and where it fails: tool choice, arguments, tool results, slow steps, cost
D. Code access	Repo access under NDA	We make the tracing change for them	Same as C, with no effort on their side

Data stays in the client's own Phoenix where possible. We sign an NDA, agree what is stored, and delete data when the engagement ends.

Pilot timeline

A first audit runs about two weeks on our side-project hours (about 10 hours each per week), with one client at a time. Re-test and monitoring follow once the client has shipped fixes.

1. **Days 1–3: Intake.** Kickoff call, intake doc, test set drafted and sent for approval.
2. **Days 3–5: Access.** Client picks an option; connection confirmed with one test chat.
3. **Days 5–8: Test and capture.** All test cases run three times; traces pulled.
4. **Days 8–11: Score and fix plan.** Checks run, human review, findings written.
5. **Days 12–14: Debrief.** Report and CSV sent, readout call held.
6. **After fixes ship: Re-test.** Same test set, before and after compared.
7. **Ongoing: Monitor.** Offered once re-runs are automated.

Quality gates

Nothing reaches a client until every box below is ticked. One wrong number in front of a client costs more trust than a late report.

- [ ] Client approved the test set before it was run.
- [ ] A person reviewed about 15% of scored cases and agreed with the scores.
- [ ] Every finding has frequency, evidence, root cause (or "needs tracing"), fix and effort.
- [ ] Every number in the report matches the CSV exactly.
- [ ] Personal data is removed from every transcript excerpt.
- [ ] Test-set version and agent version are stated, so the next run is comparable.
- [ ] The client lead can explain each finding in plain English to a non-engineer.
- [ ] Readout call booked within a week of delivery.